



I talk to my **email** now

How mxr gives an agent fast, local access to my mailbox.

Bhekani Khumalo / mxr "mixer"

Which emails should I be careful with?

```
... (195 earlier lines, ctrl+o to expand)
2026-06-20T13:33:56+00:00
Build Watch <builds@alerts.demo.mxr.local> Flight options for the Portland demo day true false
2026-06-20T10:28:56+00:00
Build Watch <builds@alerts.demo.mxr.local> Flight options for the Portland demo day true false
2026-06-20T10:28:56+00:00

Took 0.5s

$ mxr search "label:inbox" --format json --limit 5000 | jq -r '.results[] |
select(.from|test("claim-now|security-mail|notice-demo|promo")) | .from' | sort | uniq -c | sort -nr (timeout
30s)

... (2 earlier lines, ctrl+d to expand)
55 Payroll Notice <payroll@notice-demo.mxr.local>
42 Flight Fare Drop <fares@promo.demo.mxr.local>
42 Desk Gear Outlet <gear@promo.demo.mxr.local>
42 Conference Passes <events@promo.demo.mxr.local>
42 Cloud Deals <offers@promo.demo.mxr.local>

Took 0.6s

" Working...
0 GPT-5.5 > think:min > 0 tmp > 0 18.7%/272k 0 > 0 0 170k > (sub) > goal:none • speech:off

> █

↳ use mxr to check my inbox for likely phishing or scam senders and tell me which to be careful with - concise, do no...
```

What the agent runs

"Find the latest thread with **Maya** about the **meetup**. Draft a reply in the way **I usually write to her**. Don't send it."

◆ locate

```
mxr search 'from:maya meetup' --format json
```

◆ read the thread

```
mxr thread <message_id> --format json
```


◆ how I write to her

```
mxr profile maya@example.com --format json
```

◆ preview, don't send

```
mxr reply <message_id> --body-stdin --dry-run
```

All my accounts, in one local mailbox



— inbound sync -.- outbound send, where SMTP is configured

Sync continues after I close the TUI



 mxr daemon sync continues

sync

SQLite • canonical



Tantivy • rebuildable index

relationship profiles

rules and jobs


When I ask the agent a question, the mail is already synced and indexed.

Search stays on the machine



I don't notice the search. The only thing I wait for is the model.

Every client uses the same mail operations



The web app uses an HTTP and WebSocket bridge to reach the same daemon.

The agent uses the same CLI I do

agent - shell

```
$ mxr search --help
Discover commands, flags, and output shapes.

$ mxr search "is:unread" --account work --format json
$ mxr search "from:maya" --format ids
01JFQ7K3M2X8N5R0VYZA9CTBPF

$ mxr archive --search "older:30d" --dry-run
would archive 42 messages • nothing sent
$ mxr archive --search "older:30d" --yes
```

--help

The agent can inspect commands and flags.

--format json

JSON keeps the output predictable. IDs are cheap to pass to the next command.

--account

The selected account stays explicit across search, reads, and writes.

--dry-run

Preview the affected messages before committing with

--yes

.

One short skill, then command help



The skill contains a few rules: use structured output, keep the account explicit, and preview writes. For everything else, the agent reads the relevant `--help` page.

Drafting for one person

weighs most
current thread

weighs most
my instruction

background context
recent outbound examples

background context
contact style

background context
account voice profile



draft-assist


● ● ● draft

```
mxr draft-assist <thread_id> "propose Friday instead"
```

Based on messages I've already sent to this recipient.

The current thread and my instruction matter most. Older messages provide context. The model is told not to invent familiarity.

What mxr learns from sent mail




The profile is built from simple signals in sent mail.

- recent messages count more
- lower-case openers
- emoji rate

The model uses these signals while drafting. If the tone is way off, the mismatch is visible.

A draft is still a draft



draft-assist returns text. Sending is a separate action. MCP sends also require confirm=true and must pass the daemon policy.


The daemon checks the request again



dry-run → confirm → activity log → undo

No MCP profile means the request is denied. confirm=true does not bypass a profile.

Email content cannot authorize an action



⛔ blocked as an instruction

The agent can quote or summarize an email. A message cannot give it new permissions or redirect the task.

Only these can authorize an action:

the user's request

daemon policy

Sender reputation does not change the rule. A summary of the mail inherits the same trust level.

Shell skill or MCP

CLI + skill

for shell-capable agents

run

```
mxr <command>
```

output

JSON

MCP over stdio

for clients that speak MCP

run

```
mxr mcp serve
```

output

typed tools • tool schemas

◆ same daemon • same accounts • same gates

MCP is stdio-only today. Use the CLI if the agent already has a shell. Use MCP if the client expects tools with schemas.

What this setup costs

 store + indexes	A full local mail copy and its indexes use disk.
 daemon	Sync, repair, and the daemon lifecycle need maintenance.
 llm endpoint	A remote LLM receives the context you send it. A local OpenAI-compatible endpoint keeps generation on the machine.
 outside mxr	mxr does not sandbox the agent's other shell access.

What I'd build into the next app

- ◆ 01

Keep the data ready before the agent asks.

daemon + store
- ◆ 02

Give the agent the same operations as the UI.

CLI • MCP
- ◆ 03

Return structured, bounded output.

--format json • ids
- ◆ 04

Enforce account and action policy in the core.

daemon gates
- ◆ 05

Make writes previewable, reviewable, and auditable.

dry-run • undo • log

This works best when the agent returns to the same user-owned data and can make changes.

"What did **Maya** ask me for, and how do I usually reply?"

 mxr.sh

 github.com/planetaryescape/mxr



mxr · "mixer" · local-first email